

[GIT PULL] Please pull my kvm-ppc-next-4.20-1 tag

From: Paul Mackerras paulus@ozlabs.org

Paolo or Radim,

Please **do** a pull from my kvm-ppc-next-4.20-1 tag to get a PPC KVM update **for** 4.20. There is one big new feature here, which is nested virtualization support **for** radix guests on POWER9, plus a few other fixes and improvements.

A lot of the commits here are also going to be merged in the powerpc tree via the topic/ppc-kvm branch, in order to reduce conflicts between the powerpc tree and the kvm tree.

Stephen Rothwell noted that there is a conflict between my kvm-ppc-next tree and the kvm-arm tree, in that we have both allocated the next free capability number (160). It would be nice to get a stable allocation **for** the new capability that I need (one which reports/enables nested virtualization support **for** the HV KVM module) so that the necessary QEMU patches can be submitted.

Thanks,
Paul.

The following changes since commit dd5bd0a65ff6f22a32b35ca3fa1bcf7a6bc7104f:

Merge tag 'kvm-s390-next-4.20-1' of [git://git.kernel.org/pub/scm/linux/kernel/git/kvms390/linux](https://git.kernel.org/pub/scm/linux/kernel/git/kvms390/linux) into HEAD (2018-10-04 17:12:45 +0200)
are available in the git repository at:

[git://git.kernel.org/pub/scm/linux/kernel/git/paulus/powerpc](https://git.kernel.org/pub/scm/linux/kernel/git/paulus/powerpc) tags/kvm-ppc-next-4.20-1

for you to fetch changes up to 901f8c3f6feb0225c14b3bc6237850fb921d2f2d:

KVM: PPC: Book3S HV: Add NO_HASH flag to GET_SMMU_INFO ioctl result (2018-10-09 16:14:54 +1100)

PPC KVM update for 4.20.

The major new feature here is nested HV KVM support. This allows the HV KVM module to load inside a radix guest on POWER9 and run radix guests underneath it. These nested guests can run in supervisor mode and don't require any additional instructions to be emulated, unlike with PR KVM, and so performance is much better than with PR KVM, and is very close to the performance of a non-nested guest. A nested hypervisor (a guest with nested guests) can be migrated to another host and will bring all its nested guests along with it. A nested guest can also itself run guests, and so on down to any desired depth of nesting.

Apart from that there are a series of updates for IOMMU handling from Alexey Kardashevskiy, a "one VM per core" mode for HV KVM for security-paranoid applications, and a small fix for PR KVM.

Alexey Kardashevskiy (5):

- KVM: PPC: Validate all tces before updating tables
- KVM: PPC: Inform the userspace about TCE update failures
- KVM: PPC: Validate TCEs against preregistered memory page sizes
- KVM: PPC: Propagate errors to the guest when failed instead of ignoring
- KVM: PPC: Remove redundand permission bits removal

Cameron Kaiser (1):

- KVM: PPC: Book3S PR: Exiting split hack mode needs to fixup both PC and LR

Michael Ellerman (1):

Merge branch 'kvm-ppc-fixes' of paulus/powerpc into topic/ppc-kvm

Paul Mackerras (26):

KVM: PPC: Book3S HV: Provide mode where all vCPUs on a core must be the same VM

powerpc: Turn off CPU_FTR_P9_TM_HV_ASSIST in non-hypervisor mode

KVM: PPC: Book3S: Simplify external interrupt handling

KVM: PPC: Book3S HV: Remove left-over code in XICS-on-XIVE emulation

KVM: PPC: Book3S HV: Move interrupt delivery on guest entry to C code

KVM: PPC: Book3S HV: Extract PMU save/restore operations as C-callable functions

KVM: PPC: Book3S HV: Simplify real-mode interrupt handling

KVM: PPC: Book3S: Rework TM save/restore code and make it C-callable

KVM: PPC: Book3S HV: Call kvmppc_handle_exit_hv() with vcore unlocked

KVM: PPC: Book3S HV: Streamlined guest entry/exit path on P9 for radix guests

KVM: PPC: Book3S HV: Handle hypervisor instruction faults better

KVM: PPC: Book3S HV: Add a debugfs file to dump radix mappings

KVM: PPC: Use ccr field in pt_regs struct embedded in vcpu struct

KVM: PPC: Book3S HV: Use kvmppc_unmap_pte() in kvm_unmap_radix()

KVM: PPC: Book3S HV: Framework and hcall stubs for nested virtualization

KVM: PPC: Book3S HV: Nested guest entry via hypercall

KVM: PPC: Book3S HV: Use XICS hypercalls when running as a nested hypervisor

KVM: PPC: Book3S HV: Handle hypercalls correctly when nested

KVM: PPC: Book3S HV: Use hypercalls for TLB invalidation when nested

KVM: PPC: Book3S HV: Don't access HFSCR, LPIDR or LPCR when running nested

KVM: PPC: Book3S HV: Add one-reg interface to virtual PTCR register

KVM: PPC: Book3S HV: Allow HV module to load without hypervisor mode

KVM: PPC: Book3S HV: Add nested shadow page tables to debugfs

Merge remote-tracking branch 'remotes/powerpc/topic/ppc-kvm' into kvm-ppc-next

KVM: PPC: Book3S HV: Add a VM capability to enable nested virtualization

KVM: PPC: Book3S HV: Add NO_HASH flag to GET_SMMU_INFO ioctl result

Suraj Jitindar Singh (9):

KVM: PPC: Book3S HV: Clear partition table entry on vm teardown

KVM: PPC: Book3S HV: Make kvmppc_mmu_radix_xlate process/partition table agnostic
 KVM: PPC: Book3S HV: Refactor radix page fault handler
 KVM: PPC: Book3S HV: Handle page fault **for** a nested guest
 KVM: PPC: Book3S HV: Introduce rmap to track nested guest mappings
 KVM: PPC: Book3S HV: Implement H_TLB_INVALIDATE hcall
 KVM: PPC: Book3S HV: Invalidate TLB when nested vcpu moves physical cpu
 KVM: PPC: Book3S HV: Sanitise hv_regs on nested guest entry
 KVM: PPC: Book3S HV: Handle differing endianness **for** H_ENTER_NESTED

Documentation/virtual/kvm/api.txt		19	+
arch/powerpc/include/asm/asm-prototypes.h		21	+
arch/powerpc/include/asm/book3s/64/mmu-hash.h		12	+
.../powerpc/include/asm/book3s/64/tlbflush-radix.h		1	+
arch/powerpc/include/asm/hvcall.h		41	+
arch/powerpc/include/asm/kvm_asm.h		4	+-
arch/powerpc/include/asm/kvm_book3s.h		45	+-
arch/powerpc/include/asm/kvm_book3s_64.h		118	+-
arch/powerpc/include/asm/kvm_book3s_asm.h		3	+
arch/powerpc/include/asm/kvm_booke.h		4	+-
arch/powerpc/include/asm/kvm_host.h		16	+-
arch/powerpc/include/asm/kvm_ppc.h		8	+-
arch/powerpc/include/asm/ppc-opcode.h		1	+
arch/powerpc/include/asm/reg.h		2	+
arch/powerpc/include/uapi/asm/kvm.h		1	+
arch/powerpc/kernel/asm-offsets.c		5	+-
arch/powerpc/kernel/cpu_setup_power.S		4	+-
arch/powerpc/kvm/Makefile		3	+-
arch/powerpc/kvm/book3s.c		46	+-
arch/powerpc/kvm/book3s_64_mmu_hv.c		7	+-
arch/powerpc/kvm/book3s_64_mmu_radix.c		718	+++++++---
arch/powerpc/kvm/book3s_64_vio.c		89	+-
arch/powerpc/kvm/book3s_64_vio_hv.c		81	+-
arch/powerpc/kvm/book3s_emulate.c		13	+-

arch/powerpc/kvm/book3s_hv.c	864	+++++++--
arch/powerpc/kvm/book3s_hv_builtin.c	92	+--
arch/powerpc/kvm/book3s_hv_interrupts.S	95	+--
arch/powerpc/kvm/book3s_hv_nested.c	1291	+++++++
arch/powerpc/kvm/book3s_hv_ras.c	10	+
arch/powerpc/kvm/book3s_hv_rm_xics.c	13	+--
arch/powerpc/kvm/book3s_hv_rmhandlers.S	823	+++++-----
arch/powerpc/kvm/book3s_hv_tm.c	6	+--
arch/powerpc/kvm/book3s_hv_tm_builtin.c	5	+--
arch/powerpc/kvm/book3s_pr.c	5	+--
arch/powerpc/kvm/book3s_xics.c	14	+--
arch/powerpc/kvm/book3s_xive.c	63	+
arch/powerpc/kvm/book3s_xive_template.c	8	-
arch/powerpc/kvm/bookehv_interrupts.S	8	+--
arch/powerpc/kvm/emulate_loadstore.c	1	-
arch/powerpc/kvm/powerpc.c	15	+--
arch/powerpc/kvm/tm.S	250	++--
arch/powerpc/kvm/trace_book3s.h	1	-
arch/powerpc/mm/tlb-radix.c	9	+
include/uapi/linux/kvm.h	2	+
tools/perf/arch/powerpc/util/book3s_hv_exits.h	1	-

45 files changed, 3933 insertions(+), 905 deletions(-)
create mode 100644 arch/powerpc/kvm/book3s_hv_nested.c